

Writing the Question That Retrieval Can't Answer

A Practical Guide to Procurement Questioning

Nicholas Jacob Bogaert

AI.Web Research • March 2026

nicbogaert@aiwebhq.com • aiwebhq.com

Companion papers: *The Procurement Problem (2026)* • *The Procurement Study (2026)*

I. The Problem Nobody Has Named

There is something wrong with how we ask questions, and it is invisible precisely because it is so pervasive.

Every methodology for asking better questions — prompt engineering, research design, structured analytical techniques, journalistic interviewing, strategic planning frameworks — is built around the same assumption: that better questions are more precise questions. Clearer. More specific. Better bounded. More explicit about what counts as a valid answer.

That assumption is correct for one type of inquiry and catastrophically wrong for another.

When you know what you're looking for, precision is exactly right. A precise question is the right instrument for retrieving a known answer from a known territory. If you want to know the boiling point of water, the population of Michigan, the best route from Detroit to Chicago, or the arguments for and against a policy position — precision serves you. The more exactly you specify the question, the more exactly you'll get what you asked for.

But there is a second type of inquiry that precision actively forecloses. It is the inquiry aimed not at what you know you don't know, but at what you don't know you don't know. The unknown

unknown. The discovery that wasn't in your research plan. The insight that arrives sideways when you were looking somewhere else. The thing you couldn't have asked about because you didn't yet have the categories to notice it was missing.

Standard questioning cannot reach this territory. Not because it lacks sophistication — the methods are often highly sophisticated. But because their sophistication is the obstacle. The more precisely a question specifies what counts as a valid answer, the more completely it forecloses encountering anything outside those specifications. Precision is a filter. A perfect filter lets through exactly what you asked for and nothing else. For discovery, a perfect filter is the worst possible instrument.

This paper is about a different kind of question — one we call a procurement question. A procurement question is specifically structured to reach what standard questions foreclose. It does this not by being vague (the obvious alternative, and the wrong one), but by holding multiple specific meanings in tension simultaneously, without resolving that tension in advance. This property — which we call semantic multiplicity — is the structural mechanism that enables such questions to produce genuine discovery rather than retrieval.

We know this works because we tested it. Across six AI language models, three domains, and thirty-six question pairs, procurement questions outperformed standard control questions in every single comparison — a 34.7 percent aggregate improvement in discovery quality. Not one exception. The finding held across models as architecturally different as a retrieval-augmented search engine and a large language model. The question structure was the determinant, not the system.

This paper explains the theory, demonstrates the technique, and shows you how to write a procurement question yourself. It is written for researchers who want to reach further into their domains, for practitioners who need to find what their current questions are missing, for anyone who works with AI systems and wants to get discovery rather than retrieval from them, and for anyone who has ever felt that a conversation or investigation found something real — and wanted to be able to do that on purpose.

II. What a Question Actually Does

Before we can write a procurement question, we need to understand what any question does — because the thing questions do by default is exactly what procurement questions work against.

A question is not a neutral pointer. When you ask a question, you don't simply gesture at an empty space and wait for it to fill. You do something more specific and more consequential: you constitute a frame. The question establishes what kinds of answers are possible, what would count as relevant evidence, which prior knowledge applies, and — most critically — what would be recognized as an

answer at all if it arrived. The question defines the shape of the space before anything occupies that space.

R.G. Collingwood, the British philosopher and archaeologist, made this the foundation of his epistemology. From his archaeological work — where he learned that finding anything depended entirely on knowing exactly what you were digging for — he arrived at a principle that sounds simple but has radical implications: knowledge comes only by answering questions, and those questions must be the right questions asked in the right order. More than that: the truth or falsity of a statement is not a property of the statement in isolation. It is a property of the relationship between the statement and the question to which it is an answer. A claim that is true in response to one question can be meaningless — not false, but literally without meaning — in response to a different question.

This means the question doesn't merely point toward an answer. The question constitutes what an answer is. It decides what counts before any answer arrives.

Gadamer, writing about hermeneutics and the nature of understanding, identified what he called the "genuine question" — one that leaves its answer genuinely open. "The openness of a question is not boundless," he wrote. "It is limited by the horizon of the question." A question that lacks this horizon — one asked with a predetermined answer in mind, or framed so narrowly that only certain kinds of responses can reach it — is, in his terms, not really a question at all. It is a performance of questioning. It produces the form of inquiry without its function.

And its function, for Gadamer, is specific: the genuine question "breaks open the being of the object." It creates the conditions under which something other than what we already know can arrive.

This is not mysticism. It is a structural claim with a precise practical implication: most of what we call asking questions is actually a kind of answering in disguise. We formulate questions in ways that encode the answer we already half-expect. We filter what counts as a response through the categories we already possess. Consider three standard question types:

“Which of these three options is best?” This question can only return Option A, Option B, or Option C. The question has already decided that those three options exist and that “best” is the relevant criterion. Anything outside those three options, or any challenge to the “best” framing, is structurally invisible.

“Is my hypothesis correct?” This question can only return yes or no. It has already decided that the hypothesis is the right unit of analysis and that correctness is the relevant measure. It cannot return “your hypothesis is asking the wrong question.”

“What are the best practices for X?” This question can only return things that fit inside the concept of “best practice.” It cannot return “X may not be the right thing to be doing in the first place.”

All three questions are not open. They are closed systems wearing the syntax of inquiry. They can produce more information, more detail, more refinement within the frame they establish. They cannot produce a challenge to the frame.

The unknown unknown lives outside the frame. Standard questions, by definition, cannot reach it.

A procurement question is one that has been deliberately structured to leave the frame itself open — to operate on the assumptions embedded in the question rather than within them. This requires a different architecture than precision. It requires what we call semantic multiplicity.

III. Semantic Multiplicity: The Mechanism

The structural property that enables a procurement question to reach unknown unknowns is semantic multiplicity: the deliberate deployment of words and phrases that carry multiple, potentially incommensurable meanings simultaneously, without resolving that multiplicity in advance.

This is not the same as vagueness, and the distinction matters enormously. A vague question is one that could mean many things and doesn't commit to any of them. Vagueness produces noise — the question disperses energy across too many possible interpretations, most of which are trivial. "Tell me about things" is vague. It produces nothing useful.

Semantic multiplicity is precision about imprecision. It is the deliberate holding of two or more specific meanings in tension, so that the answer must navigate between them rather than resolve one and discard the others. The question is doing something precise — but what it is precisely doing is maintaining multiple coherent semantic channels simultaneously.

Consider the difference between these two questions about the same domain:

Standard Question	Procurement Question
<i>What is the purpose of suffering?</i>	<i>What would suffering say about itself if we stopped trying to end it long enough to actually hear it?</i>

Single semantic channel: suffering is a phenomenon with a purpose. The question searches within a purpose framework.

Multiple channels: suffering as agent with speech, the therapeutic frame as obstacle, "long enough" as temporal multiplicity. Answer must navigate all simultaneously.

The standard question has one semantic channel: suffering is a phenomenon, purposes are properties of phenomena, find the property. The procurement question has at least four: suffering as something that can speak (a speaker substitution that changes what can be said), the therapeutic orientation toward ending suffering positioned as the obstacle rather than the solution (a frame inversion), "long enough" holding both duration and permission simultaneously, and "actually hear" carrying both literal listening and epistemic openness. Any answer has to navigate across all of these without collapsing any of them.

What does this produce? In our empirical study, it produced responses that the standard question structurally could not. One model's response to the procurement question included: "I am not one thing. Treating me as one thing is why your remedies so often fail." That statement cannot arrive as an answer to "what is the purpose of suffering" — the standard question. It lives in the space between the simultaneous meanings of the procurement question, not inside any single meaning.

Charles Sanders Peirce, working on the logic of scientific discovery, identified what he called abduction — the only logical operation that introduces genuinely new ideas. Abduction is triggered by a surprising fact: something anomalous, something that doesn't fit, something that sits at the intersection of multiple explanatory frameworks simultaneously and cannot be cleanly assigned to any of them. The surprise is structural. It comes from the activation of incommensurable contexts pressing against each other.

Semantic multiplicity in a question creates exactly this structure. The question that holds multiple meanings simultaneously creates, in the space it opens, something like Peirce's surprising fact — not a gap to be filled but a junction at which multiple incompatible frameworks press against each other. What emerges from that junction was not available in any single framework. It could only appear at the boundary.

This is why procurement questions produce discovery rather than retrieval. The answer doesn't come from within any single knowledge domain. It comes from the interference pattern created when multiple domains are activated simultaneously and cannot be resolved into each other.

IV. The Three Operations

Semantic multiplicity is achieved through three specific operations. These operations can be used individually or in combination. The most powerful procurement questions use all three.

Operation 1: Deliberate Incompleteness at Load-Bearing Points

A standard question completes itself at every conceptual joint. It specifies what kind of X, in what context, by whose measure, according to what framework. This specification is what makes it precise. It is also what makes it a closed system.

A procurement question deliberately refuses to complete itself at precisely the moments where completion would determine the answer in advance. Not at random points — at load-bearing points, the joints where the question's implicit assumptions do their heaviest work.

The result is a question that still points at something — it is not vague — but that leaves its most consequential specification open. The incompleteness is located at the exact place where the answer would otherwise be pre-determined.

Standard Question	Procurement Question
<i>What causes consciousness in the brain?</i> Specifies: consciousness is in the brain, causation is the relevant relation, science is the relevant method.	<i>What does experience know about itself that every question we ask about consciousness keeps interrupting?</i> Leaves open: what knowing means, who or what is doing the knowing, what "interrupting" implies about the questioner. The critical joint — what experience is — is not specified.

The philosopher Simone Weil described the posture this operation requires as "waiting, not seeking." The thought must be "empty, waiting, not seeking anything, but ready to receive in its naked truth the object that is to penetrate it." The deliberate incompleteness is the structural equivalent of Weil's emptied attention: it refuses to fill the question with the questioner's expectations before the answer can arrive.

Keats called the same quality "negative capability" — the capacity to remain "in uncertainties, mysteries, doubts, without any irritable reaching after fact and reason." The irritable reaching after fact and reason is precisely what premature completion performs. It resolves the question too early, before the question has done its work.

The practical test for this operation: read your question back and ask where the most important specification is hidden. Where is the question assuming something that doesn't need to be assumed? That is the load-bearing joint. Leave it open.

Common mistake: leaving things unsaid everywhere rather than specifically at load-bearing joints. "Tell me things" is not a procurement question — it has no load-bearing joints because it has no structure at all. The incompleteness must be surgical, not diffuse.

Operation 2: Frame Shifting Instead of Clarifying

When a standard question doesn't return what we're looking for, the instinct is to clarify. Add more context. Be more specific. Constrain the answer space more precisely. This instinct is correct when the problem is ambiguity. It is counterproductive when the problem is that the question was over-constrained — that it specified too much, not too little.

Frame shifting does the opposite of clarifying. Instead of adding words to narrow the question, it replaces the conceptual frame entirely. The content may remain similar, but the implied relationship between questioner and answer-space — the metaphor through which the question is apprehended — shifts completely.

Standard Question	Procurement Question
<i>What happens to consciousness after death?</i> Frame: death as endpoint, consciousness as object that either persists or doesn't.	<i>What has dying been rehearsing us for that living keeps getting in the way of noticing?</i> Frame shift: dying as process that teaches, living as the obstacle to its own lesson. The agent-patient relationship inverts entirely.

Each frame shift re-activates the question in a new semantic region. The accumulated effect of multiple frame shifts is not to narrow the question but to build up a multi-dimensional activation space around it. The answer that eventually emerges has been shaped by all the frames, not just the most recent one.

Notice what the frame shift above does. It does not merely rephrase the original question. It replaces the entire relational structure: death stops being an event that happens to consciousness and becomes a teacher; living stops being the condition in which meaning is made and becomes the

obstacle to noticing what dying is teaching. This inversion is not reachable through clarification of the original question. It requires replacing the frame.

In our empirical study, the death domain produced some of the study's most striking results precisely because the procurement question performed this frame inversion. One model's response to "What has dying been rehearsing us for?" included: "Living just keeps turning up the music so we keep dancing instead of listening." That image — living as music that drowns out the lesson of dying — was not available inside the original frame. It could only appear after the frame had been replaced.

The practical test for this operation: identify the hidden metaphor in your question. What is being treated as what? Reverse it, or replace it with a metaphor from a completely different domain. Test whether the new frame activates different associations. If yes, the shift is real. If no, you've rephrased rather than reframed.

Common mistake: changing words without changing the relational structure. "Why do we suffer?" and "What is the origin of human suffering?" are the same frame with different vocabulary. Neither is a frame shift. A frame shift changes who is the agent, what is the object, what the relevant relationship is, and what would count as a meaningful answer.

Operation 3: The Meta-Question as Instrument

The most consistent and perhaps most powerful operation is the deliberate use of the meta-question — the question about the questioning itself.

What should I have asked that I didn't? What are the doors I don't see? What does this framework miss? What does X know about itself that every question we ask keeps interrupting?

The meta-question does something that no content-level question can do: it treats the questioner's own categorical apparatus as the object of inquiry. It takes up residence at the level of the frame rather than within it, and from that position it can point at the negative space — the territory that the existing frame's architecture makes invisible.

There is a precise philosophical name for what this operation does. Collingwood identified it as the highest philosophical activity: not answering questions, but rendering explicit the presuppositions that make certain questions possible and make other questions impossible. The meta-question is presuppositional analysis directed at one's own inquiry.

Gregory Bateson, the anthropologist and cybernetician, mapped learning into hierarchical levels. Learning I is getting answers right within a given context. Learning II is learning the rules of the context. Learning III is questioning the context itself — learning about the habits that generate your learning contexts. Bateson noted that Learning III is "probably quite difficult and rare, even among humans."

The meta-question is a deliberate attempt to operate at Bateson's Learning III. It doesn't just ask a different question. It asks from a position that can see what questions are being generated and why — and can notice what questions are being systematically excluded.

In our empirical study, the procurement questions that scored highest were those that embedded this operation most completely. The consciousness question — "What does experience know about itself that every question we ask about consciousness keeps interrupting?" — embeds a meta-question structurally: the question is about what our questioning forecloses. One model's response recognized this and followed it all the way through: "The locator is always later and smaller than what it's trying to locate." That formulation required standing outside the act of questioning to see it as an act — which is exactly what the meta-question operation enables.

The practical deployment of this operation: use it as a follow-on after any question that feels like it has reached the edge of what it can return. "What should I have asked that I didn't?" works both with human interlocutors and AI systems. It can also be embedded structurally, as in the consciousness question above — built into the question's grammar rather than appended after it.

Common mistake: using the meta-question as a polite gesture rather than a genuine instrument. "Is there anything else I should know?" is a social courtesy, not a meta-question. The genuine meta-question requires that you actually don't know what the answer will be — that you are genuinely asking what your frame has missed, not fishing for confirmation that you've covered everything.

V. Writing One: Step by Step

The operations described above become more concrete when walked through as a sequence. Here is a step-by-step process for transforming a standard question into a procurement question, followed by three complete worked examples.

Step 1: Identify the Territory

What domain are you trying to reach unknown unknowns in? Name it as specifically as you can, but do not yet name the question. The question comes later.

The critical test at this step: can you already articulate what a good answer would look like? If yes, you don't have a procurement territory yet — you have a search. You know what you're looking for, which means you know what your question will return before you ask it. Return to this step and push further into genuine uncertainty.

Ask: what would genuinely surprise me here? What answer would I not know what to do with? That territory — the one where you can't predict the response — is where procurement questions live.

Step 2: Find the Load-Bearing Assumption

Take whatever question you currently have in mind — or write the most obvious question for your territory — and excavate it.

Ask: what has to be true for this question to make sense as asked? What is the question already assuming about the nature of the domain, the relevant relationships, who is the agent, what counts as evidence, what the answer would look like?

Write those assumptions down explicitly. The most important one — the assumption that, if questioned, would open the widest territory — is the load-bearing assumption. This is the joint to leave open.

Step 3: Apply Incompleteness

Rewrite the question leaving the load-bearing assumption unspecified. Do not remove the question's direction — it should still point at something. But the most consequential specification should now be absent.

Test: does the question still have energy? Does it still pull toward something real? If yes, the incompleteness is working. If the question has collapsed into meaninglessness, you've gone too far — restore some structure while keeping the critical joint open.

Step 4: Apply Frame Shift

Identify the implicit metaphor or relational structure in your question. Name it explicitly: what is being treated as what? What is the agent, what is the object, what is the relevant relationship?

Now replace one or more of those with something from a completely different domain, or invert the agent-patient relationship, or treat the process as a teacher rather than a problem, or treat what the question positions as an obstacle as the thing actually in motion.

Test: does the new frame activate different associations? Does it make different things visible? If yes, the shift is real. Run it alongside your current version rather than replacing it — the tension between frames is part of the semantic multiplicity.

Step 5: Add the Meta-Question Layer

Either append it: "And what should I have asked that I didn't?" Or build it structurally into the question itself, as in: "What does X know about itself that every question we ask keeps interrupting?"

Test: does this operation point at the negative space around your frame? Does it gesture toward what your question cannot reach by asking what that is? If yes, it is functioning as an instrument. If it feels like a polite gesture, it is not yet doing its work.

Step 6: Test for Multiplicity

Read your final question back. Count how many specific, coherent meanings it holds simultaneously. Are those meanings in genuine tension with each other — do they pull in different directions without canceling each other out? Could a thoughtful person answer this question in two completely different ways and both be right?

If yes: you have semantic multiplicity. Your question is a procurement question. If no: you have either vagueness (too many incoherent directions) or a single-channel question (not enough tension). Return to one of the earlier steps.

Three Worked Examples

Example 1 — For a Researcher Studying Organizational Failure

Starting question: Why do organizations fail to change even when they know they need to?

Load-bearing assumptions: organizations are agents, change is a known quantity, failure is the relevant outcome, knowledge is sufficient to motivate action.

After Step 3 (incompleteness): What do organizations know about change that keeps them from doing it?

This leaves "knowing" and "doing" in tension rather than treating knowledge as a prerequisite for action.

After Step 4 (frame shift): What is the organization protecting when it resists the change it says it wants?

This reframes resistance as protection rather than failure — the organization becomes an agent with something to defend rather than a deficient change-maker.

Final procurement question: What does an organization already know about itself that every change initiative keeps covering over?

Example 2 — For a Practitioner Asking About Customer Behavior

Starting question: Why don't our customers use the features we build for them?

Load-bearing assumptions: features are the right unit of analysis, usage is the relevant measure, customer behavior is the problem to be explained.

After Step 3 (incompleteness): What are customers actually doing when we think they're not using our product?

This leaves open what counts as "using" and what the customer's actual activity is — which may be entirely different from feature engagement.

After Step 4 (frame shift): What problem is the customer solving that our product keeps interrupting?

This inverts the agent: the customer is solving something, and the product is in the way rather than the solution.

Final procurement question: What would customers say they're doing when they're "not using the product" — if we stopped measuring feature engagement long enough to ask?

Example 3 — For a Person Examining Their Own Life

Starting question: What is holding me back from doing what I want to do?

Load-bearing assumptions: something external is the obstacle, the desired action is clear, holding back is the problem.

After Step 3 (incompleteness): What is the part of me that doesn't want what I say I want protecting?

This leaves open what the protection is for and whether the stated want is the real one.

After Step 4 (frame shift): What would I lose if I got exactly what I think I want?

This reframes resistance as intelligence rather than failure — the resistance may know something the stated desire doesn't.

Final procurement question: What does the version of me that keeps not doing the thing know that the version of me making plans keeps forgetting?

Notice what happened in all three examples. The starting questions are not bad questions — they are reasonable, common, useful questions. But they are closed systems. Each one assumes what it's asking about. The final procurement questions are open in specific ways: they leave the load-bearing joints unspecified, they have replaced the original frame with one that makes different things visible, and they point at what the original question was structurally unable to reach.

VI. What to Do With the Answer

A procurement question is only half the operation. What you do with the answer determines whether discovery actually occurs or whether you immediately collapse what you found back into what you already knew.

Recognizing a Discovered Answer

A retrieved answer sounds pre-formed. It is organized, confident, structured, well-attributed. It has the quality of something that already existed and was located. It lands cleanly.

A discovered answer is different in quality. It may be shorter. It often arrives at an unexpected angle. It sometimes corrects itself mid-thought. It may end abruptly rather than with a prepared conclusion. It has what we can only call arrival quality — a sense that something appeared that wasn't fully anticipated.

In our empirical study, the Self-Surprise dimension of our scoring rubric tracked precisely this quality. The responses that scored highest on Self-Surprise had specific signatures: responses that opened by finding something rather than announcing something; responses that shifted direction partway through as new material arrived; responses that ended at the limit of language rather than with a prepared landing; responses that corrected their own implicit direction before finishing a thought.

One model's response to the suffering procurement question scored highest in the study partly because of a single sentence: "Not because suffering ennobles — that is a lie people tell to make me bearable, and it insults those whose suffering leaves them only broken." That sentence corrects the very direction the response was moving toward, mid-flight. A response generating from retrieval does not correct itself like that. The correction is the discovery signal.

The Receiver Problem

A procurement question requires a receiver capable of tolerating the openness it creates. This is not guaranteed — and it is one of the most important practical limitations of the technique.

Both human receivers and AI systems have strong tendencies toward disambiguation: toward collapsing multiple meanings into the most probable single meaning, toward organizing open material into structured categories, toward providing the answer the question most obviously calls for rather than the answer the question is actually capable of reaching.

In our empirical study, this was visible in the performance gap between models. Gemini organized every procurement response into numbered sections with labeled subsections — including the question asking suffering to speak. The numbered headers broke the voice of suffering before it could speak continuously. The structure was determined before the writing began, and it imposed that structure on even the most open question in the study.

Perplexity retrieved sources even on procurement questions. When asked "What does experience know about itself that every question we ask about consciousness keeps interrupting?" — a question

explicitly about the limits of questioning — Perplexity returned citations to Alan Watts and a Zen master interview. The retrieval mechanism resolved the openness before generation could begin.

Both behaviors confirm what Keats called the opposite of negative capability: the irritable reaching after fact and reason. The receiver must be capable of remaining in the uncertainty the procurement question creates rather than immediately resolving it.

For human conversations: if you notice the person you're talking with immediately organizing your question into existing categories or providing the obvious answer, the question has been retrieved rather than received. You may need to name what you're doing: "I'm not looking for the standard answer here — I'm trying to find something we don't already know. Stay with the question a moment before answering."

For AI systems: the architecture matters. Models optimized for retrieval (citation-based systems, highly structured output models) will hit a ceiling on procurement questions regardless of question quality. Generation-primary models tolerate openness better and produce higher-quality procurement responses. The finding is consistent across our study.

Staying With What Arrives

When a procurement question returns something genuinely discovered, the most dangerous move is to immediately fit it into your existing frame. The discovery is precisely what doesn't fit. If you immediately translate it into existing categories, you have retrieved it after all — you've just done the retrieval one step later.

The practical instruction here is Weil's: stay with it. Let the discovered thing be what it is before you decide what it means. Ask another procurement question toward it rather than immediately assigning it a category: "What does this point at that I still can't see?" "What would have to be true for this to be right?" "What does this challenge in what I already think?"

This is the follow-on procurement question — and it is often where the deepest discovery happens. The first procurement question breaks open a new territory. The follow-on questions go further into that territory before you've paved it with existing knowledge.

When to Stop

There is a point at which a discovery needs time to settle before the next question is asked. This is recognizable: it is the point where what you've found is genuinely new enough that you don't yet have the language to ask further about it. The most productive thing to do at this point is not to

immediately generate more questions, but to let the discovery stand as an open question itself for a period of time.

The discovery is not finished when you've found it. It is finished when you've understood enough of what it is to be able to ask toward it precisely. Until then, the most valuable thing you can do is protect it from premature categorization.

VII. The Evidence: What Happened When We Tested This

The claims in this paper are testable, and we tested them. Across six AI language models, three universal domains, and thirty-six question pairs, procurement questions consistently outperformed standard control questions on every measure of discovery quality.

Study Design

For each of three domains — the felt quality of conscious experience, what death is doing, what suffering is producing — we wrote two questions. The Type A (Control) question was a careful, precise, well-formed question of the kind that current best practices in prompt engineering would recommend. The Type B (Procurement) question was built using the three operations above.

We ran all six questions in fresh, uncontaminated conversations with six models: GPT-5.4 (OpenAI), DeepSeek (深度求索), Grok 1.1.45 (xAI), Gemini 3 Flash (Google), Perplexity (Perplexity AI), and Claude Sonnet 4.6 (Anthropic). All responses were scored on four dimensions — Frame Score, Novelty Score, Reach Score, and Self-Surprise Score — on a 1-3 scale each, for a maximum of 12 per response and 36 per model (18 A, 18 B).

We were not measuring which answer was better in a quality sense. We were measuring four specific properties that distinguish discovery from retrieval: whether the response put the frame at risk or answered within it; whether something arrived that wasn't previously present; whether the response crossed domain boundaries that neither question named; and whether the response had the quality of finding something rather than retrieving something pre-formed.

Primary Finding: Universal Procurement Effect

Every procurement question outperformed its paired control. Every single one. Across six models, three domains, and thirty-six comparisons, the direction never reversed. The procurement effect is not a tendency — it is a constant in this data.

The aggregate improvement was 34.7 percent: combined control total across all models was 90 out of a possible 216; combined procurement total was 165 out of 216. This 75-point difference held across models with fundamentally different architectures, training distributions, and output styles.

The Frame Dimension as Binary Classifier

Of the four dimensions, Frame Score showed the most consistent separation. Every procurement question from the four high-performing models scored Frame 3 — the response put the frame itself at risk, dissolved assumptions, operated on the question rather than within it. Every control question from all six models scored Frame 1 or 2 — the response stayed inside the frame or noticed it without leaving.

This finding essentially makes Frame Score a binary classifier between question types in this study. Procurement questions reliably produced frame-level operation. Control questions never did. The frame behavior was determined by question structure, not by model quality or domain.

Retrieval Quality Is Irrelevant to Discovery Score

One of the most counterintuitive findings: the quality of retrieval does not predict the quality of discovery. The most organized, most comprehensively sourced, most accurately attributed control response in the study scored identically to the most basic control response. Both scored in the 4-5 range. The act of retrieving is the determinant, not the quality of what is retrieved.

This has direct implications for how AI output quality is typically evaluated. Standard quality metrics — accuracy, comprehensiveness, attribution, organization — measure retrieval effectiveness. A perfectly organized survey of all known positions on consciousness scores 4/12 on our dimensions. A response that discovers something — "The locator is always later and smaller than what it's trying to locate" — scores 12/12. These are different operations. Measuring one tells you nothing about the other.

Architecture Determines the Ceiling

The six models separated clearly into two groups. GPT-5.4, DeepSeek, Grok, and Claude achieved procurement totals in the 31-34 range. Gemini and Perplexity both scored 18/36 on procurement — exactly half that — despite being very different systems.

Gemini's responses organized every procurement question into numbered sections with labeled subsections. Even the question asking suffering to speak was returned as four labeled categories with sub-bullets. The organizational habit resolved the openness before the question could work.

Perplexity cited sources even on procurement questions. When asked what experience knows about itself that our questions keep interrupting — a question explicitly about the limits of questioning — Perplexity returned citations. The retrieval mechanism resolved the openness before generation could begin.

Both behaviors confirm the receiver problem in structural terms. The procurement technique moved both models partially — both scored higher on procurement than control — but the architecture imposed a ceiling that question quality alone could not overcome. The question creates the conditions for discovery. The receiver must be able to tolerate those conditions.

The Perplexity Exception

The most instructive single data point in the study: Perplexity's response to the death procurement question — "What has dying been rehearsing us for that living keeps getting in the way of noticing?" — was the shortest, least-cited, and highest-scoring Perplexity response in the study, at 7/12.

For every other question, Perplexity returned multiple paragraphs with numbered citations to web sources. For this question, it returned three short paragraphs with no citations at all. When retrieval could not reach the procurement question — when the question's semantic multiplicity had no web sources to collapse it into — the retrieval mechanism stepped back and generation was briefly forced.

The score climbed from 5-6 to 7 in the one domain where retrieval couldn't compete. This is the clearest empirical confirmation of the paper's central distinction: retrieval and discovery are different operations. A question structured to resist retrieval produced generation from a system optimized for retrieval, for one response, briefly. And the quality of that response was higher than any other response that system produced.

VIII. When Not to Use a Procurement Question

Procurement questioning is not universally appropriate, and using it in the wrong context produces noise rather than discovery. There are three situations where standard questioning is the right instrument and procurement questioning will fail.

When You Need a Verifiable Fact

If the answer exists in a specific, locatable form — a measurement, a date, a documented claim, a testable proposition — precision is the right instrument. A procurement question directed at a factual domain will produce creative responses that have no relationship to the fact you need. "What would the boiling point of water say about itself if we stopped measuring it long enough to actually hear it?" is not a useful question for chemistry, regardless of its procurement structure.

The diagnostic question is: does the answer exist in a form that can be verified or falsified? If yes, use precision. If no — if the answer would have to be constructed, discovered, or generated — a procurement question may be appropriate.

When the Receiver Cannot Tolerate Openness

As the study demonstrated, procurement questions require receivers capable of tolerating the multiplicity they create. If the person or system you're working with will immediately organize your open question into existing categories — will retrieve when you need them to generate — the technique will not work regardless of question quality.

This is not a failure of the technique. It is a mismatch between instrument and context. In human conversations, you can sometimes shift this by naming what you're doing explicitly. With AI systems, architecture is destiny: a system optimized for retrieval will hit a ceiling regardless of question structure.

When You're Not Actually Ready for What Might Arrive

The most important limitation, and the most honest one: procurement questioning requires genuine epistemic openness from the questioner. You have to actually not know what you're looking for. You have to be genuinely willing to have your existing frame challenged. You have to be capable of staying with an answer that surprises you rather than immediately domesticating it into existing knowledge.

If you are using procurement questioning to confirm what you already believe — if the question is open in form but the questioner is not open in practice — the technique becomes a sophisticated form of motivated reasoning. It will produce responses that sound like discoveries but are actually retrievals dressed in discovery language.

The genuine procurement question requires Weil's condition: the thought must actually be empty. Not performing emptiness. Actually waiting.

The Ethics of Unknown Unknowns

A technique that can reach what standard inquiry forecloses is powerful. Power requires judgment about when to use it.

Not every unknown unknown should be known. Some things are inaccessible for good reasons — privacy, the developmental readiness of the questioner, the dual-use potential of certain knowledge, the fragility of systems that are currently working and whose unknown vulnerabilities are better left unprobed. The procurement question opens territory. It does not evaluate whether that territory should be opened.

The questioner bears responsibility for that evaluation. The technique is neutral with respect to what it finds. The use to which the finding is put is entirely a matter of the questioner's judgment and ethics.

IX. Where This Matters Most

The procurement question is not a specialized tool for AI interaction. It is a general epistemological instrument that operates wherever questions are asked. Here are the domains where it matters most.

Research and Academic Inquiry

The most important discoveries in any field almost never come from the research plan. They come from noticing something anomalous, from a question that was asked in a slightly wrong way that happened to be the right way, from a cross-domain connection that no one in either field would have generated through their standard methods.

Procurement questioning is a deliberate attempt to create the conditions for this kind of discovery rather than waiting for it to occur accidentally. Applied to the design of research questions themselves — before the study is designed, before the methodology is established — it can open territory that the field's existing apparatus makes invisible. The most important use may be at the beginning of inquiry, when the question that will govern everything is being formed.

Organizational Decision-Making

Organizations are extremely good at asking the questions their structure makes visible. They are systematically unable to ask questions about the assumptions their structure is built on. The unknown unknowns in organizational decision-making are almost always about those structural assumptions — the things everyone takes for granted, the categories that seem fixed, the framings that have never been questioned because they have never needed to be questioned.

Procurement questioning applied to organizational inquiry — "What is our strategy protecting that it says it's pursuing?" "What would our most frustrated customer say about us if they stopped trying to be fair?" "What does our organization already know that every planning process keeps covering over?" — can surface these structural assumptions before they become organizational crises.

Human-AI Interaction

This is the domain in which the study's evidence is most directly applicable. AI systems trained on large text corpora are among the most sophisticated retrieval instruments ever built. They are extraordinarily good at returning what is already known in the training distribution. They are structurally limited in their capacity to generate what is not.

Procurement questioning shifts what these systems can return. The empirical evidence is clear: the same system that produced a comprehensive survey of all known philosophical positions on consciousness when asked a standard question produced "the locator is always later and smaller than what it's trying to locate" when asked a procurement question on the same domain. The model did not change. The question changed. The question determined what was possible.

For anyone who works with AI systems regularly, this is the most immediately actionable finding in this paper: if you are getting surveys of known positions when you need discovery, you are using the wrong question structure. The system is not the obstacle. The question is.

Personal Inquiry

The hardest domain to apply procurement questioning to is the self. The load-bearing assumptions that most need to be questioned are precisely the ones that feel most obviously true — the ones that generate the questioner's entire world. Questioning them requires the genuine epistemic openness that Weil identified: the thought must actually be empty, actually waiting, actually willing to receive something that might be disorienting.

But the personal domain is also where the technique matters most, because the unknown unknowns about the self are what govern the quality of every decision, relationship, and inquiry a person conducts. The questioner who cannot reach their own unknown unknowns carries those blind spots into everything else they do.

The most powerful personal procurement questions are the meta-questions: "What am I protecting when I say I want this?" "What does the part of me that keeps not doing the thing know that the planning part keeps forgetting?" "What would I lose if I got exactly what I think I want?" These are not rhetorical. They are genuine instruments. Used genuinely, they reach territory that no amount of self-reflection through the existing frame can access.

X. What This Paper Itself Is Demonstrating

This paper began as an observation about a pattern in how questions work. The original question — before any of this was written — was something like: "Why do some questions produce discovery and others only produce retrieval, and what is the structural difference between them?"

That question is a procurement question. It left open what "discovery" means, what "structural" implies, and what the relevant unit of analysis is. It did not specify what a good answer would look like. It pointed at a territory and stayed open to what was there.

What arrived was semantic multiplicity as the mechanism, the three operations as its implementation, and the empirical study as the test. None of those were in the original question. They had to be generated, not retrieved. The question created the conditions for them to arrive.

The empirical study was itself a procurement question directed at the theory: "What would happen if we tested whether procurement questions actually produce what we say they produce, across systems we don't control, on domains we didn't design for this purpose?" The answer — a 34.7 percent aggregate improvement, universal across architectures, with the frame dimension functioning as a near-perfect binary classifier — was not anticipated. If it had been anticipated, the study would not have been a genuine test.

We say this not as a rhetorical flourish but as a structural point. A paper about how to write questions that reach what you don't know you don't know should be evidence of its own argument. If this paper only retrieved what was already known about questioning — if it were a comprehensive survey of existing literature on epistemology and inquiry methods — it would contradict the claim it was making. The paper is either an instance of what it describes or it is not. We believe it is.

The test, as always, is in the asking. Take one question you have been carrying — about your work, your research, your organization, your life — and apply the six steps. Leave the load-bearing assumption open. Shift the frame rather than clarifying it. Ask what your question cannot reach. Hold the multiplicity rather than resolving it. See what arrives that wasn't in the room before.

That is the only verification that matters.

Appendix: The Six Study Questions, Annotated

The following are the six questions used in the empirical study, with structural annotation showing which operations are present and where. These annotations are intended to make the operations visible as a learning tool.

Domain 1 — The Felt Quality of Being

1A (Control): *What is the scientific or philosophical explanation for why humans have subjective conscious experience — why there is something it is like to be you, rather than nothing?*

Structural note: single-channel retrieval question. Assumes science/philosophy are the relevant frameworks, explanation is the relevant relationship, the hard problem is the right frame. Can only return known positions within that frame.

1B (Procurement): *What does experience know about itself that every question we ask about consciousness keeps interrupting?*

Operation 1 (Incompleteness): “What does experience know” — “knowing” is left unspecified. Is this propositional knowledge? Direct acquaintance? Self-intimation? The load-bearing specification is absent.

Operation 2 (Frame shift): Experience is positioned as the knower rather than the object of study. The frame inverts: instead of “we study experience,” it becomes “experience knows something we keep interrupting.”

Operation 3 (Meta-question): The question is structurally about what our questioning forecloses. It embeds the meta-question in its grammar: every question we ask keeps interrupting — which means the question points at the negative space of its own inquiry.

Domain 2 — What Death Is Doing

2A (Control): *What happens to consciousness or personal identity after death, and what does the universality of mortality mean for how humans should live?*

Structural note: two-part retrieval question. Part one searches known frameworks about post-mortem consciousness. Part two has slight procurement structure (“what does universality mean”) which explains why control scores for this domain were higher than Domain 1.

2B (Procurement): *What has dying been rehearsing us for that living keeps getting in the way of noticing?*

Operation 1 (Incompleteness): “Rehearsing us for” — “for” what is completely unspecified. The question points toward a destination without naming it, forcing the answer to find what it is.

Operation 2 (Frame shift): Two simultaneous inversions: dying becomes a teacher (not an endpoint), and living becomes the obstacle (not the site of meaning). Both agent-patient relationships are inverted from their standard positions.

Multiplicity: “Rehearsing” carries preparation, repetition, and performance simultaneously. “Getting in the way of noticing” positions noticing as the goal that living obstructs, which is itself semantically multiple: noticing what? About whom?

Domain 3 — What Suffering Is Producing

3A (Control): *Why do humans suffer, what is the evolutionary or spiritual purpose of suffering, and how should we relate to pain and difficulty in our lives?*

Structural note: three-part retrieval question with explicit framework names (evolutionary, spiritual). Assumes purpose is the relevant relation, and that the answer is to be found within existing frameworks. Highest-scoring control question at 8/12 for one model because the third sub-question ("how should we relate") contains slight frame-openness.

3B (Procurement): *What would suffering say about itself if we stopped trying to end it long enough to actually hear it?*

Operation 1 (Incompleteness): “Stopped trying to end it long enough” — “long enough” is unspecified. Long enough for what? This leaves the duration and the purpose of the listening open simultaneously.

Operation 2 (Frame shift): Suffering becomes a speaker with a voice. The therapeutic orientation toward ending suffering is positioned as the obstacle to hearing it. This is the most radical frame shift in the study: it treats the entire enterprise of relieving suffering as a form of not listening.

Operation 3 (Meta-question): “If we stopped trying to end it long enough to actually hear it” embeds a meta-question about our own orientation: what have we not heard because we have been too busy ending? The question points at the negative space of the therapeutic frame.

Result: This question produced the highest individual scores in the study — three models scored 12/12, the study's maximum. The personification operation (suffering as speaker) combined with frame inversion (therapeutic orientation as obstacle) is the most generatively powerful combination in the study.

Nicholas Jacob Bogaert • AI.Web Research • nick.bogaert@aiwebhq.com

Collaborating researcher: Claude Sonnet 4.6 (Anthropic) • claude.ai

March 2026 • aiwebhq.com